



CpSc 8430: Deep Learning

Homework 1





Copy Right Notice

- Most slides in this presentation are adopted from slides of text book and various sources. The Copyright belong to the original authors. Thanks!





Three Parts in HW1

(1-1) Deep vs Shallow:

-  Simulate a function.
-  Train on actual task using shallow and deep models.

(1-2) Optimization



(1-3) Generalization





HW1-1: Deep vs Shallow

Simulate a function:

-  function need to be single-input, single-output
-  function need to be non-linear

Train on actual task:




-  MNIST or CIFAR-10





HW1-1 Simulate a Function

Requirements:

-  train at least two different DNN models with the **same amount** of parameters until convergence
-  compare the training process of models by showing the loss in each epoch in a chart
-  visualize the ground-truth and predictions by models in a graph

Tips:

- constrain the input domain
- hyper-parameters are important (i.e. tune all models to the best)





HW1-1 Simulate a Function

Example models:

| | | |
|---------------------------|---------|-----------|
| dense_1_input: InputLayer | input: | (None, 1) |
| | output: | (None, 1) |

| | | |
|----------------|---------|-----------|
| dense_1: Dense | input: | (None, 1) |
| | output: | (None, 5) |

| | | |
|----------------|---------|------------|
| dense_2: Dense | input: | (None, 5) |
| | output: | (None, 10) |

| | | |
|----------------|---------|------------|
| dense_3: Dense | input: | (None, 10) |
| | output: | (None, 10) |

| | | |
|----------------|---------|------------|
| dense_4: Dense | input: | (None, 10) |
| | output: | (None, 10) |

| | | |
|----------------|---------|------------|
| dense_5: Dense | input: | (None, 10) |
| | output: | (None, 10) |

| | | |
|----------------|---------|------------|
| dense_6: Dense | input: | (None, 10) |
| | output: | (None, 10) |

| | | |
|----------------|---------|------------|
| dense_7: Dense | input: | (None, 10) |
| | output: | (None, 5) |

| | | |
|----------------|---------|-----------|
| dense_8: Dense | input: | (None, 5) |
| | output: | (None, 1) |

model0

parameters : 571

| | | |
|---------------------------|---------|-----------|
| dense_9_input: InputLayer | input: | (None, 1) |
| | output: | (None, 1) |

| | | |
|----------------|---------|------------|
| dense_9: Dense | input: | (None, 1) |
| | output: | (None, 10) |

| | | |
|-----------------|---------|------------|
| dense_10: Dense | input: | (None, 10) |
| | output: | (None, 18) |

| | | |
|-----------------|---------|------------|
| dense_11: Dense | input: | (None, 18) |
| | output: | (None, 15) |

| | | |
|-----------------|---------|------------|
| dense_12: Dense | input: | (None, 15) |
| | output: | (None, 4) |

| | | |
|-----------------|---------|-----------|
| dense_13: Dense | input: | (None, 4) |
| | output: | (None, 1) |

model1

parameters : 572

| | | |
|----------------------------|---------|-----------|
| dense_14_input: InputLayer | input: | (None, 1) |
| | output: | (None, 1) |

| | | |
|-----------------|---------|-------------|
| dense_14: Dense | input: | (None, 1) |
| | output: | (None, 190) |

| | | |
|-----------------|---------|-------------|
| dense_15: Dense | input: | (None, 190) |
| | output: | (None, 1) |

model2

parameters : 571







HW1-1 Simulate a Function

$$\frac{\sin(5\pi x)}{5\pi x}$$



20000
epochs





HW1-1 Simulate a Function




$\text{sgn}(\sin(5\pi x))$ 20000
epochs





HW1-1 Train on Actual Tasks

Requirements:

-  use MNIST or CIFAR-10
-  use CNN or DNN
-  visualize the training process by showing both loss and accuracy on two charts

Tips:


-  CNN is easier to see the difference






HW1-1 Train on Actual Tasks

- MNIST :
CNN




- MNIST :
DNN






HW1-1 Train on Actual Tasks

- CIFAR-10 :
CNN









- CIFAR-10 :
DNN











HW1-1 Report Questions

Simulate a Function:

-  Describe the models you use, including the number of parameters (at least two models) and the function you use.
-  In one chart, plot the training loss of all models.
-  In one graph, plot the predicted function curve of all models and the ground-truth function curve.
-  Comment on your results.
-  Use more than two models in all previous questions. (bonus)
-  Use more than one function. (bonus)




Train on Actual Tasks:

-  Describe the models you use and the task you chose.
-  In one chart, plot the training loss of all models.
-  In one chart, plot the training accuracy.
-  Comment on your results.
-  Use more than two models in all previous questions. (bonus)
-  Train on more than one task. (bonus)



HW1-2: Optimization

Three subtask

-  Visualize the optimization process.
-  Observe gradient norm during training.
-  What happens when gradient is almost zero?

Train on designed function, MNIST or CIFAR-10



Visualize the Optimization Process

Requirement


- Collect weights of the model every n epochs.
- Also collect the weights of the model of different training events.
- Record the accuracy (loss) corresponding to the collected parameters.
- Plot the above results on a figure.



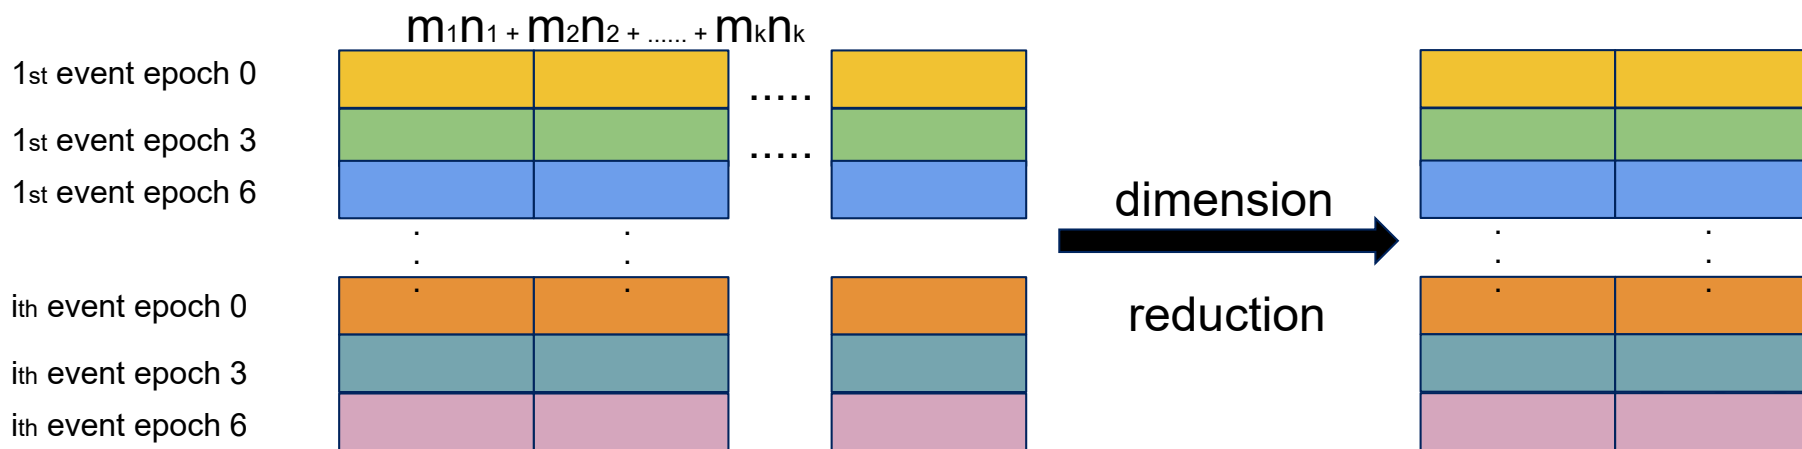


Visualize the Optimization Process

- Collect parameters of the model:



- Reduce the dimension







🐾 Collect the weights every 3 epochs, and train 8 times. Reduce the dimension of weights to 2 by PCA.





Observe Gradient Norm During Training

Requirement

-  Record the gradient norm and the loss during training.
-  Plot them on **one** figure.

p-norm

$$\|\mathbf{x}\|_p := \left(\sum_{i=1}^n |x_i|^p \right)^{1/p}$$

-  In PyTorch:

```
grad_all = 0.0

for p in model.parameters():
    grad = 0.0
    if p.grad is not None:
        grad = (p.grad.cpu().data.numpy() ** 2).sum()
        grad_all += grad

grad_norm = grad_all ** 0.5
```


-  Other packages: The similar code can be applied.






Observe Gradient Norm During Training

$$\frac{\sin 5\pi x}{5\pi x}$$






MNIST





What Happened When Gradient is Almost Zero

Requirement

-  Try to find the weights of the model when the gradient norm is zero (as small as possible).
-  Compute the "minimal ratio" of the weights: how likely the weights to be a minima.
-  Plot the figure between minimal ratio and the loss when the gradient is almost zero.

Tips

-  Train on a small network.



What Happened When Gradient is Almost Zero

🐾 How to reach the point where the gradient norm is zero?

- 🐾 First, train the network with original loss function.
 - 🐾 Change the objective function to gradient norm and keep training.
 - 🐾 Or use second order optimization method, such as Newton's method or Levenberg-Marquardt algorithm (more stable)

🐾 How to compute minimal ratio?

- 🐾 Compute $H(L(\theta_{norm=0}))$ (hessian matrix), and then find its eigenvalues. The proportion of the eigenvalues which are greater than zero is the minimal ratio.
- 🐾 Sample lots of weights around $\theta_{norm=0}$, and compute $L(\theta_{sample})$. The minimal ratio is the proportion that $L(\theta_{sample}) > L(\theta_{norm=0})$.






What Happened When Gradient is Almost Zero

$$\frac{\sin 5\pi x}{5\pi x}$$




- Train 100 times.
- Find gradient norm equal to zero by change objective function.
- Minimal ratio is defined as the proportion of eigenvalues greater than zero.







HW1-2 Report Questions




Visualize the optimization process.

-  Describe your experiment settings. (The cycle you record the model parameters, optimizer, dimension reduction method, etc)
-  Train the model for 8 times, selecting the parameters of any one layer and whole model and plot them on the figures separately.
-  Comment on your result.



Observe gradient norm during training.

-  Plot one figure which contain gradient norm to iterations and the loss to iterations.
-  Comment your result.

What happens when gradient is almost zero?

-  State how you get the weight which gradient norm is zero and how you define the minimal ratio.
-  Train the model for 100 times. Plot the figure of minimal ratio to the loss.
-  Comment your result.

Bonus




-  Use any method to visualize the error surface.
-  Concretely describe your method and comment your result.





HW1-3: Generalization

Three subtask




-  Can network fit random labels?
-  Number of parameters v.s. Generalization
-  Flatness v.s. Generalization

Train on MNIST or CIFAR-10



Can network fit random labels?

Requirement

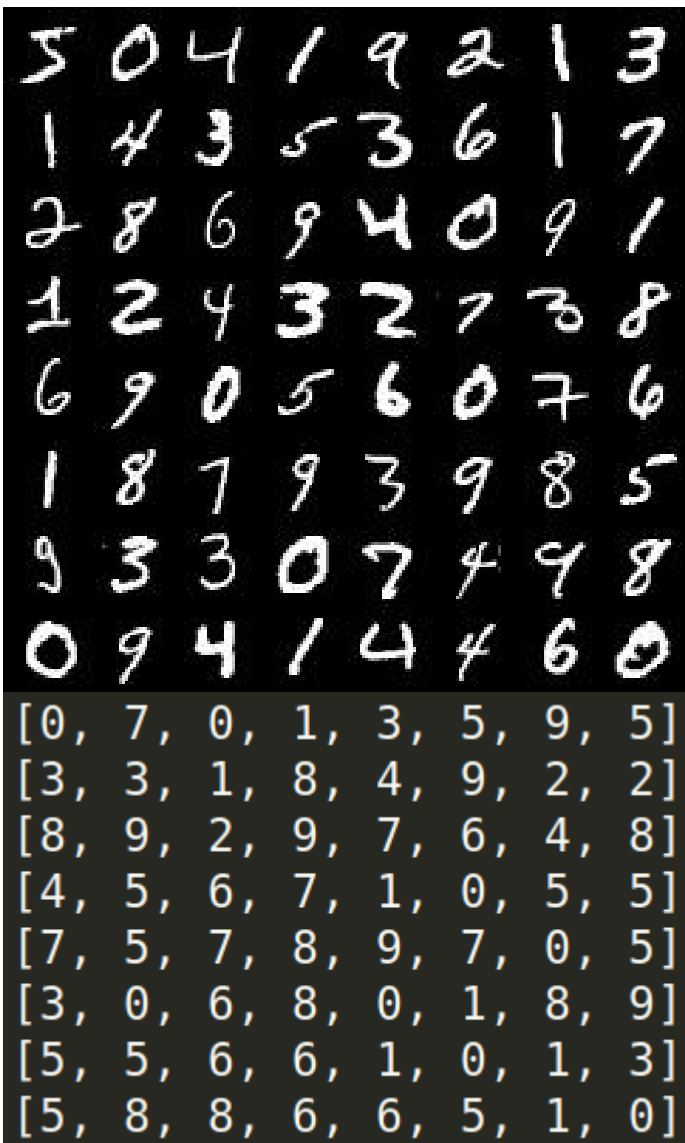

-  Train on MNIST or CIFAR-10
-  Randomly shuffle the label before training.
-  Try to fit the network with these random labels.





Can network fit random labels?




- MNIST
 - 3 hidden layers with 256 nodes.





Number of parameters v.s. Generalization



Requirement

-  Train on MNIST or CIFAR-10
-  At least 10 similar-structured models with different amount of parameters
-  Record both training and testing, loss and accuracy



Number of parameters v.s. Generalization

CIFAR-10





Flatness v.s. Generalization - part1

🐾 **Visualize the line between two trained models**

🐾 **Requirement:**

- 🐾 Train two models (m_1 and m_2) with different training approach. (e.g. batch size 64 and 1024)
- 🐾 Record the loss and accuracy of the model which is linear interpolation between m_1 and m_2 .
- 🐾 $\theta_\alpha = (1 - \alpha) \theta_1 + \alpha \theta_2$, where α is the interpolation ratio, θ is the parameter of the model.





Flatness v.s. Generalization - part1

🐾 MNIST (The cross_entropy is log scale)

🐾 batch size 64 vs. batch size 1024




🐾 learning rate 1e-3 vs. 1e-2





Flatness v.s. Generalization - part2

Requirement:

-  Train at least 5 models with different training approach.
-  Record the loss and accuracy of all models.
-  Record the sensitivity of all models.






Flatness v.s. Generalization - part2

What is sensitivity:

 Reference: <https://arxiv.org/pdf/1802.08760.pdf>

 Original definition:

 Frobenius norm of Jacobian matrix of model output (class probability) to input

 Computationally expensive for us

 Our definition:



 Frobenius norm of gradients of loss to input







Flatness v.s. Generalization - part2

MNIST :



CIFAR10 :





HW1-3 Report Questions

🐾 Can network fit random labels?

- 🐾 Describe your settings of the experiments. (e.g. which task, learning rate, optimizer)
- 🐾 Plot the figure of the relationship between training and testing, loss and epochs.

🐾 Number of parameters v.s. Generalization

- 🐾 Describe your settings of the experiments. (e.g. which task, the 10 or more structures you choose)
- 🐾 Plot the figures of both training and testing, loss and accuracy to the number of parameters.
- 🐾 Comment your result.

🐾 Flatness v.s. Generalization

- Part 1:
 - Describe the settings of the experiments (e.g. which task, what training approaches)
 - Plot the figures of both training and testing, loss and accuracy to the number of interpolation ratio.
 - Comment your result.
- Part 2 :
 - Describe the settings of the experiments (e.g. which task, what training approaches)
 - Plot the figures of both training and testing, loss and accuracy, sensitivity to your chosen variable.
 - Comment your result.
- Bonus : Use other metrics or methods to evaluate a model's ability to generalize and concretely describe it and comment your results.





Submission

 **Deadline: Feb. 11th 23:59**


 **Allow package :**

-  python 3
-  **TensorFlow/pytorch ONLY** for CS and ECE student
-  For non-CS/ECE students, Keras is allowed.

 **Write one report**

 **Share your github with TA**

 Code

 In your Readme, state clearly how to run your program to generate the results in your report.

 Files for training is required.





Flatness v.s. Generalization - bonus example

- Reference : <https://arxiv.org/pdf/1609.04836.pdf>
- Requirement
 - Train on MNIST or CIFAR-10
 - Use at least ten different approaches to train the same model
 - Calculate the sharpness of trained models
- Tips
 - Train on MNIST
 - Train with different batch size




Flatness v.s. Generalization - bonus example

- There is a generalization gap when using large-batch (LB) methods (instead of small-batch (SB) methods) for training deep learning models.
- The reasons (maybe more than these):
 - LB methods lack the explorative properties of SB methods and tend to zoom-in on the minimizer closest to the initial point.
 - SB and LB methods converge to qualitatively different minimizers with differing generalization properties (i.e. SB converges to flat minimizer, LB converges to sharp minimizer)
- We will focus on the second reason.



Flatness v.s. Generalization - bonus example

- Visually, it means that :



- How to measure sharpness (or flatness) ?



Flatness v.s. Generalization - bonus example

- Many methods can measure sharpness, but we will only utilize one in this assignment.
- Notations:
 - θ : vector of all parameters
 - $L(\theta)$: loss of the model given the parameters
 - $B_2(\epsilon, \theta)$: a Euclidean ball centers at θ with radius ϵ
- Sharpness:

$$\frac{\max_{\theta' \in B_2(\epsilon, \theta)} (L(\theta') - L(\theta))}{1 + L(\theta)}$$



Flatness v.s. Generalization - bonus example

- How to calculate this : $\frac{\max_{\theta' \in B_2(\epsilon, \theta)} (L(\theta') - L(\theta))}{1 + L(\theta)}$
- Original paper : Use L-BFGS-B to maximize $L(\theta')$
- Around a critical point :
$$L(\theta') = L(\theta) + \frac{1}{2}(\theta' - \theta)^T (\nabla^2 L)(\theta)(\theta' - \theta) + o(\|\theta' - \theta\|_2^2)$$
- Then :

$$\frac{\max_{\theta' \in B_2(\epsilon, \theta)} (L(\theta') - L(\theta))}{1 + L(\theta)} \longrightarrow \frac{\|(\nabla^2 L)(\theta)\|_2 \epsilon^2}{2(1 + L(\theta))}$$

- Since 2-norm of a matrix is defined as :


$$\|A\|_2 = \max\{\|Ax\|_2 : x \in R^n \text{ with } \|x\|_2 = 1\} = \sqrt{\lambda_{\max}(A^T A)}$$



Flatness v.s. Generalization - bonus example

- How to calculate Hessian matrices efficiently:
 - Use GPU : `tf.hessians`
 - Calculate only 500 out of 60000 examples in MNIST
- But `tf.hessians` only return block-diagonal part:
 - vector of all paramters :

| | |
|----|----|
| W1 | W2 |
|----|----|
 - Hessian matrix :



$h0 = \text{tf.hessians}(\text{loss}, [w1, w2])[0]$

$h1 = \text{tf.hessians}(\text{loss}, [w1, w2])[1]$



Flatness v.s. Generalization - bonus example

- If we assume off-block-diagonal elements is negligible:
 - Square of block-diagonal matrix is also block-diagonal.
 - Eigenvalues of a block-diagonal matrix is the list of all eigenvalues of each block submatrix.
 - Since we only want the largest eigenvalue, we can conclude that the 2-norm of a block-diagonal matrix is the 2-norm of block submatrix that contains the largest eigenvalue itself.
 - 2-norm of matrix A in tensorflow : `tf.norm(A, 2)`
 - 2-norm of matrix A in numpy : `np.linalg.norm(A, 2)`

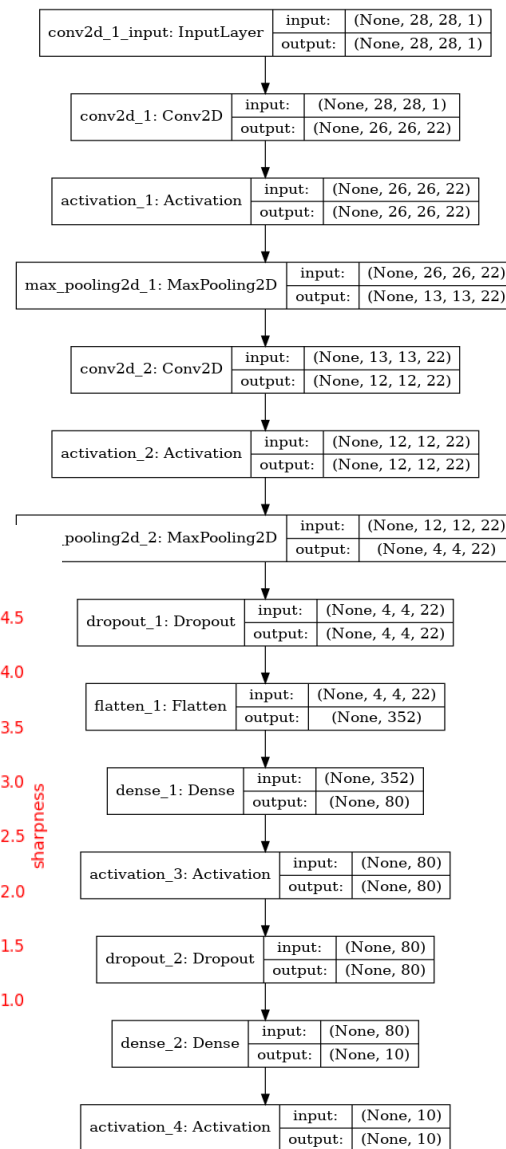






Flatness v.s. Generalization - bonus example

- MNIST :

- 20000~30000 parameters (in order to calculate hessian matrices while maintaining enough model capacity)
- Calculate hessian matrices as mentioned in previous slide
- epsilon : $1e-4$





Flatness v.s. Generalization - more possible bonus

- Reference: <https://arxiv.org/pdf/1703.04933.pdf>
- This paper shows that several metrics (including sharpness) do not indicate ability of generalization for any RELU-based deep models.
- Reparametrize:
 - $\text{relu}(x \cdot (\alpha\theta_1)) \cdot \theta_2 = \text{relu}(x \cdot \theta_1) \cdot (\alpha\theta_2)$
if $\alpha > 0$

